

# Convolutional Networks

Motivation, working with images, trainable kernels

Machine Learning and Data Mining, 2020

Artem Maevskiy

National Research University Higher School of Economics



LAMBA • HSE

November 13, 2020

# How to work with image-like data?




# Working with images

- ▶ Extremely high-dimensional input
  - E.g. even a small 640x480 color image would make up almost 1M input features (pixel brightness levels in R, G and B)
  - So a fully-connected hidden representation with just 100 units would require 100M parameters


# Working with images

- ▶ Extremely high-dimensional input
  - E.g. even a small 640x480 color image would make up almost 1M input features (pixel brightness levels in R, G and B)
  - So a fully-connected hidden representation with just 100 units would require 100M parameters
- ▶ Is quite data-hungry to train when using fully-connected layers:
  - Identifying an object on a picture would require examples with all possible locations of that object on the picture



# Translational symmetry


- ▶ A cat moved from one part of an image to another is still a cat
- ▶ Why don't we use the same model to look at different patches of an image trying to identify the object of interest:



Image




Object of interest




Probability of locating the object

# Translational symmetry


- ▶ A cat moved from one part of an image to another is still a cat
- ▶ Why don't we use the same model to look at different patches of an image trying to identify the object of interest:



Image




Object of interest




Probability of locating the object

# Translational symmetry


- ▶ A cat moved from one part of an image to another is still a cat
- ▶ Why don't we use the same model to look at different patches of an image trying to identify the object of interest:



Image




Object of interest




Probability of locating the object

# Translational symmetry


- ▶ A cat moved from one part of an image to another is still a cat
- ▶ Why don't we use the same model to look at different patches of an image trying to identify the object of interest:



Image



Object of interest




Probability of locating the object




# Translational symmetry


- ▶ A cat moved from one part of an image to another is still a cat
- ▶ Why don't we use the same model to look at different patches of an image trying to identify the object of interest:



Image




Object of interest




Probability of locating the object

# Translational symmetry


- ▶ A cat moved from one part of an image to another is still a cat
- ▶ Why don't we use the same model to look at different patches of an image trying to identify the object of interest:



Image




Object of interest




Probability of locating the object

# Translational symmetry

- ▶ A cat moved from one part of an image to another is still a cat
- ▶ Why don't we use the same model to look at different patches of an image trying to identify the object of interest:



Image



Object of interest

<b>0.001</b>	<b>0.003</b>	<b>0.876</b>
<b>0.000</b>	<b>0.002</b>	<b>0.001</b>

Probability of locating the object

# Translational symmetry

- ▶ A cat moved from one part of an image to another is still a cat
- ▶ Why don't we use the same model to look at different patches of an image trying to identify the object of interest:

**This may be implemented with a 2D convolution!**





Image  
**Input**



Object of interest  
**Convolutional  
kernel**

**0.001    0.003    0.876**

**0.000    0.002    0.001**

Probability of locating the object  
**Output**

# 2D convolution



# 2D convolution

$$\text{Output}(i, j) = \sum_{i', j'} \text{Input}(i', j') \cdot \text{Kernel}(i' - i, j' - j)$$

<i>a</i>	<i>b</i>	<i>c</i>
<i>d</i>	<i>e</i>	<i>f</i>
<i>g</i>	<i>h</i>	<i>i</i>

Input

<i>x</i>	<i>y</i>
<i>z</i>	<i>w</i>

Kernel

$ax + by + dz + ew$	

Output

# 2D convolution

$$\text{Output}(i, j) = \sum_{i', j'} \text{Input}(i', j') \cdot \text{Kernel}(i' - i, j' - j)$$

$a$	$b$	$c$
$d$	$e$	$f$
$g$	$h$	$i$

Input

$x$	$y$
$z$	$w$

Kernel

$ax + by + dz + ew$	$bx + cy + ez + fw$

Output

# 2D convolution

$$\text{Output}(i, j) = \sum_{i', j'} \text{Input}(i', j') \cdot \text{Kernel}(i' - i, j' - j)$$

$a$	$b$	$c$
$d$	$e$	$f$
$g$	$h$	$i$

Input

$x$	$y$
$z$	$w$

Kernel

$ax + by + dz + ew$	$bx + cy + ez + fw$
$dx + ey + gz + hw$	

Output



# 2D convolution

$$\text{Output}(i, j) = \sum_{i', j'} \text{Input}(i', j') \cdot \text{Kernel}(i' - i, j' - j)$$

$a$	$b$	$c$
$d$	$e$	$f$
$g$	$h$	$i$

Input

$x$	$y$
$z$	$w$

Kernel

$ax + by + dz + ew$	$bx + cy + ez + fw$
$dx + ey + gz + hw$	$ex + fy + hz + iw$

Output

# 2D convolution

$$\text{Output}(i, j) = \sum_{i', j'} \text{Input}(i', j') \cdot \text{Kernel}(i' - i, j' - j)$$

$a$	$b$	$c$
$d$	$e$	$f$
$g$	$h$	$i$

Input

$x$	$y$
$z$	$w$

Kernel

$ax + by$ $+ dz + ew$	$bx + cy$ $+ ez + fw$
$dx + ey$ $+ gz + hw$	$ex + fy$ $+ hz + iw$

Output

- Different kernels may extract different features

# Examples



Input

► Blur:

$$\text{kernel} = \begin{pmatrix} 1 & 4 & 7 & 4 & 1 \\ 4 & 16 & 26 & 16 & 4 \\ 7 & 26 & 41 & 26 & 7 \\ 4 & 16 & 26 & 16 & 4 \\ 1 & 4 & 7 & 4 & 1 \end{pmatrix} \cdot \frac{1}{273}$$



Output

# Examples



Input

- Sharpen:

$$\text{kernel} = \begin{pmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{pmatrix}$$



Ouput

# Examples



Input

- Edge detection:

$$\text{kernel} = \begin{pmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{pmatrix}$$



Output


# Examples



Input

- ▶ Edge detection:

$$\text{kernel} = \begin{pmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{pmatrix}$$




Ouput

- ▶ In the context of deep learning, the kernel parameters are **trainable**
- ▶ I.e. the network **learns the kernel** to **extract useful features**

# 2D convolution as a matrix multiplication

- ▶ Unwrap the 2D images into 1D vectors
- ▶ Re-write the convolution as a regular matrix-vector multiplication
- ▶ I.e. fully-connected layers comprise convolutions
  - Yet they are much more complex



$ax + by + dz + ew$
$dx + ey + gz + hw$
$bx + cy + ez + fw$
$ex + fy + hz + iw$

=

$x$	$y$	0	$z$	$w$	0	0	0	0
0	$x$	$y$	0	$z$	$w$	0	0	0
0	0	0	$x$	$y$	0	$z$	$w$	0
0	0	0	0	$x$	$y$	0	$z$	$w$

×

$a$
$b$
$c$
$d$
$e$
$f$
$g$
$h$
$i$


# 2D convolutional layers





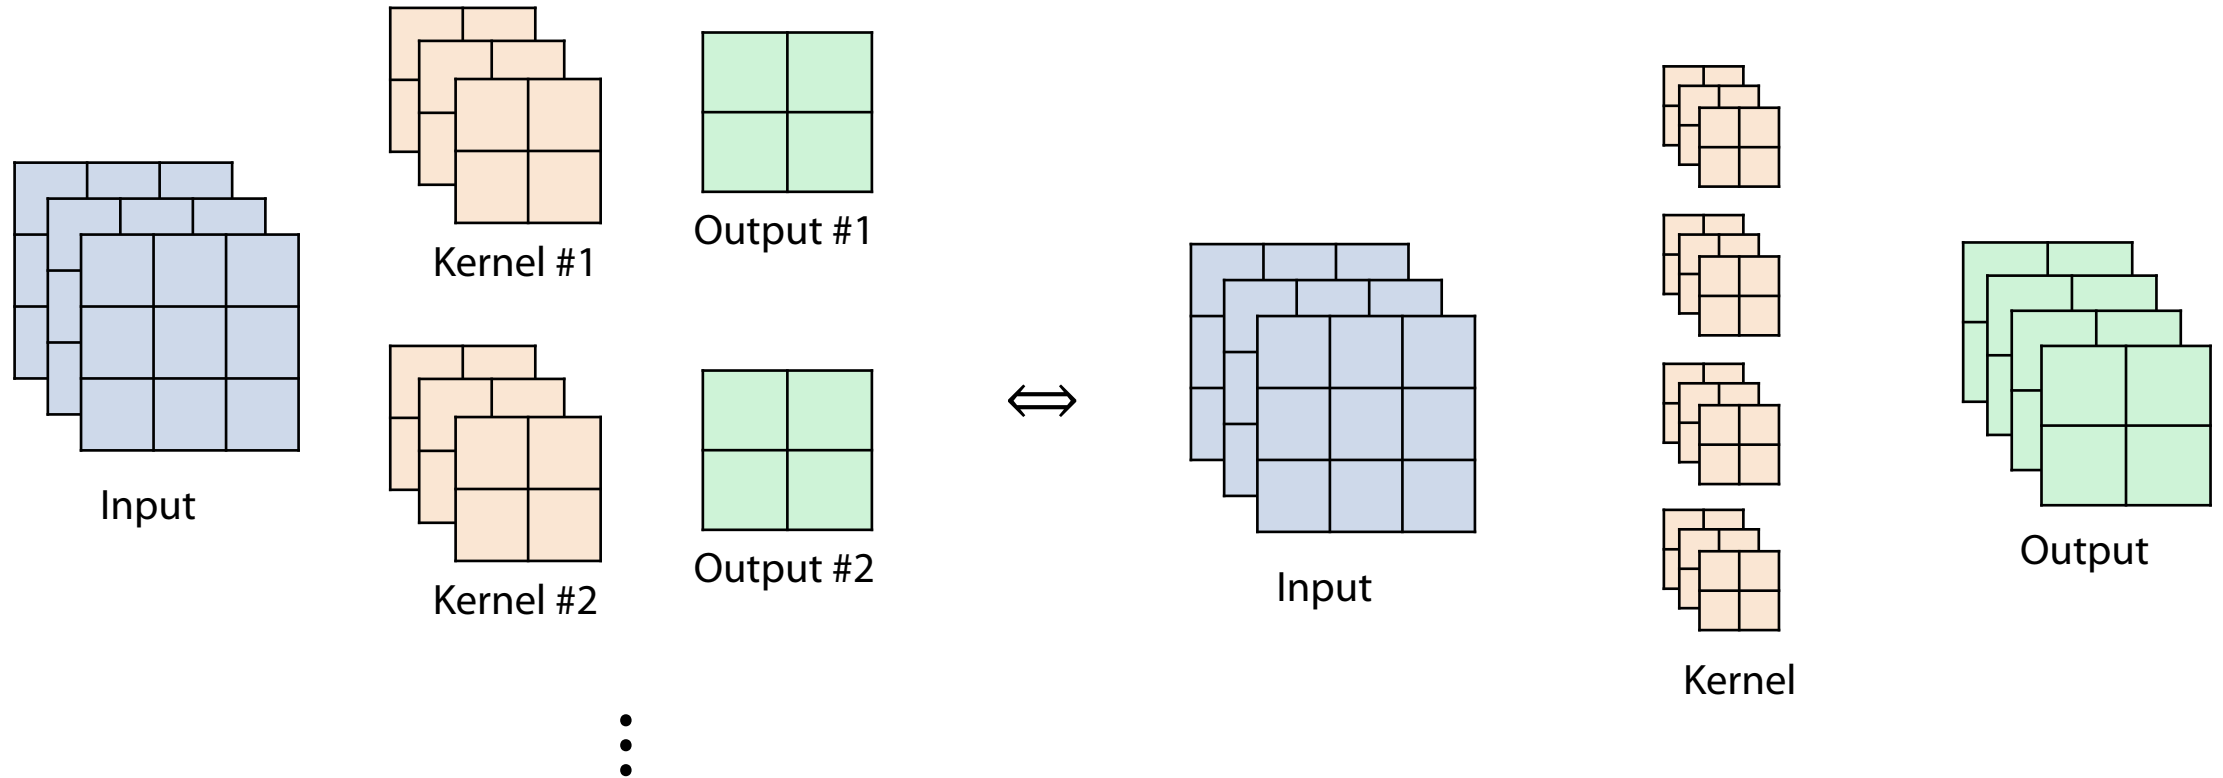
# Input channels

- ▶ In practice images have multiple channels
  - E.g. 3 color channels of a color image



# Output channels

- ▶ In practice we want to extract multiple features




- ▶ Kernel becomes 4D:  $H_K \times W_K \times C_{in} \times C_{out}$
- ▶ Output becomes 3D:  $H \times W \times C_{out}$

# Putting it all together

$$\text{Output}(i, j, c_{out}) = \sum_{i', j', c_{in}} \text{Input}(i', j', c_{in}) \cdot \text{Kernel}(i' - i, j' - j, c_{in}, c_{out}) + \overset{\text{bias term}}{\downarrow} b_{out}$$

- ▶ Note: with this approach the output width/height is smaller than the input width/height
  - By how much (for a given kernel width and height)?
- ▶ Sometimes the border of the input image is **padding** with some values (e.g. s.t. the output has the same size)
  - Controlled by the “padding” parameter



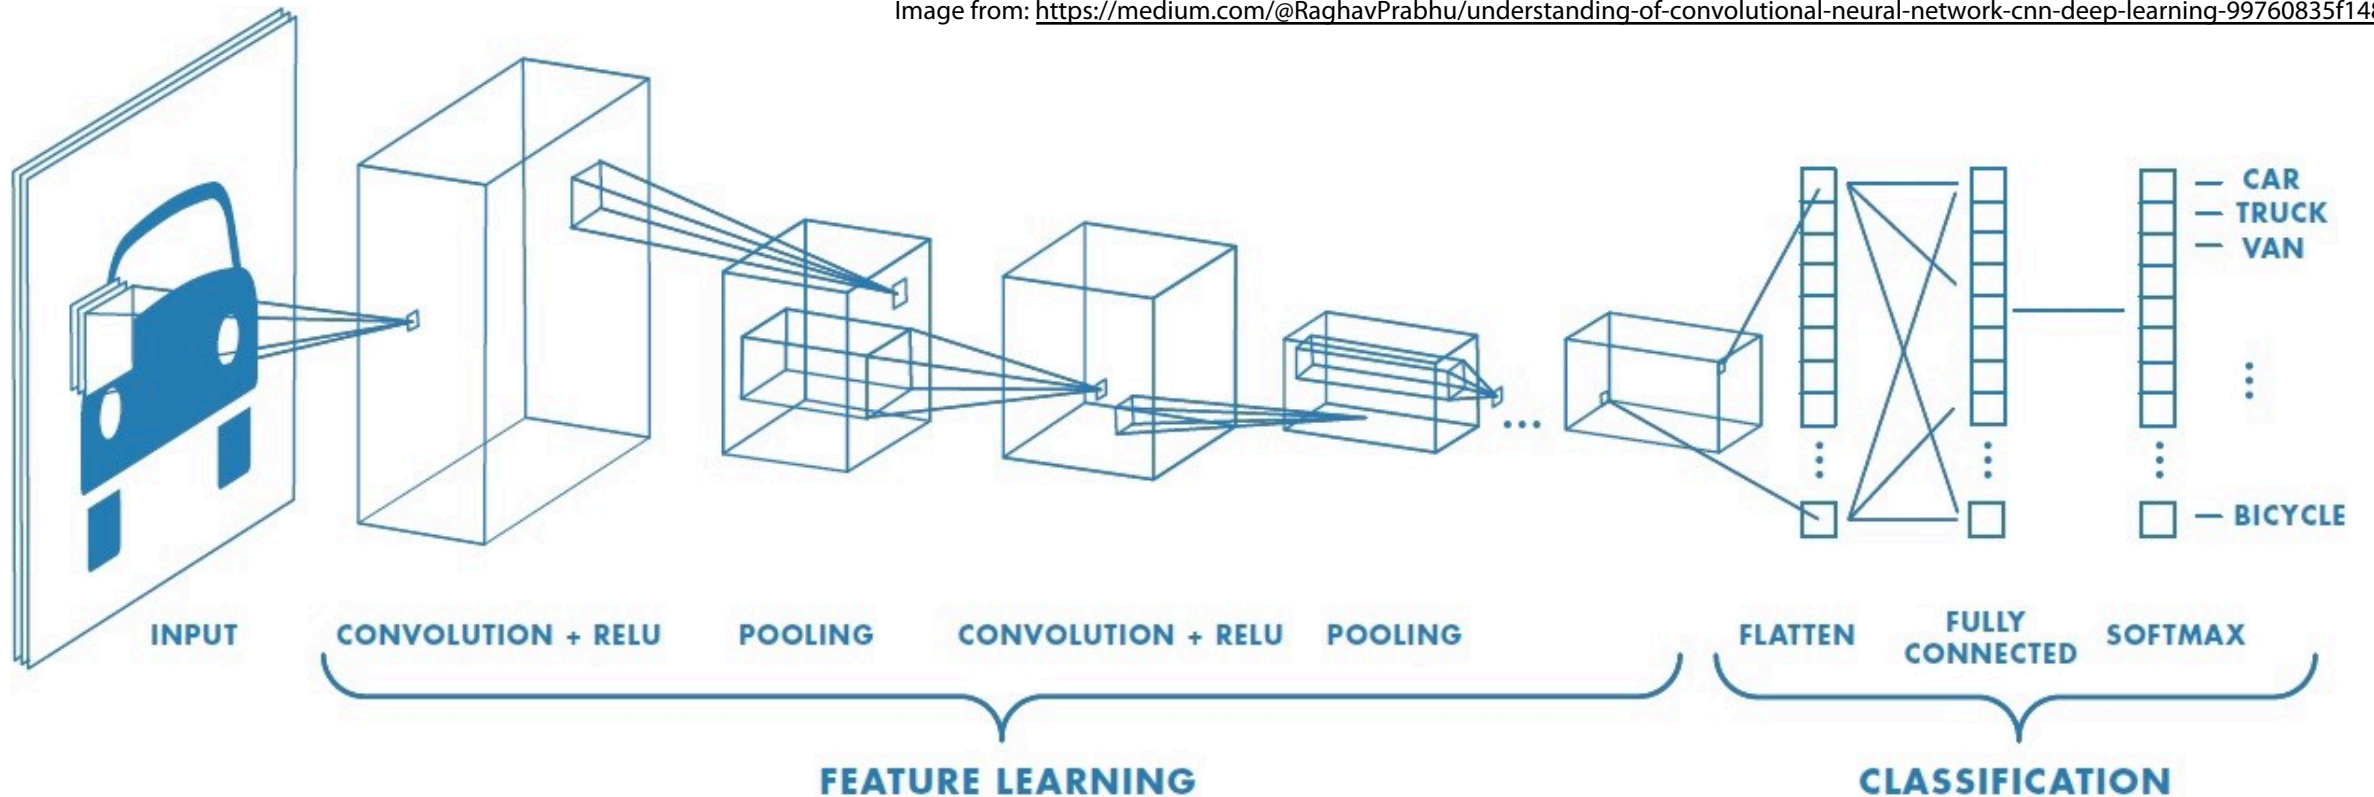
- ▶ Some other parameters:
  - “stride” – by how many pixels the kernel window steps (equals 1 in the examples here)
  - “dilation” – kernel “spread” (e.g. see [this animation](#))

# Typical network architecture



# Deep convolutional network

Image from: <https://medium.com/@RaghavPrabhu/understanding-of-convolutional-neural-network-cnn-deep-learning-99760835f148>



- Convolution is typically followed by a pooling operation (aggregating values of nearby pixels), e.g. maxpool:

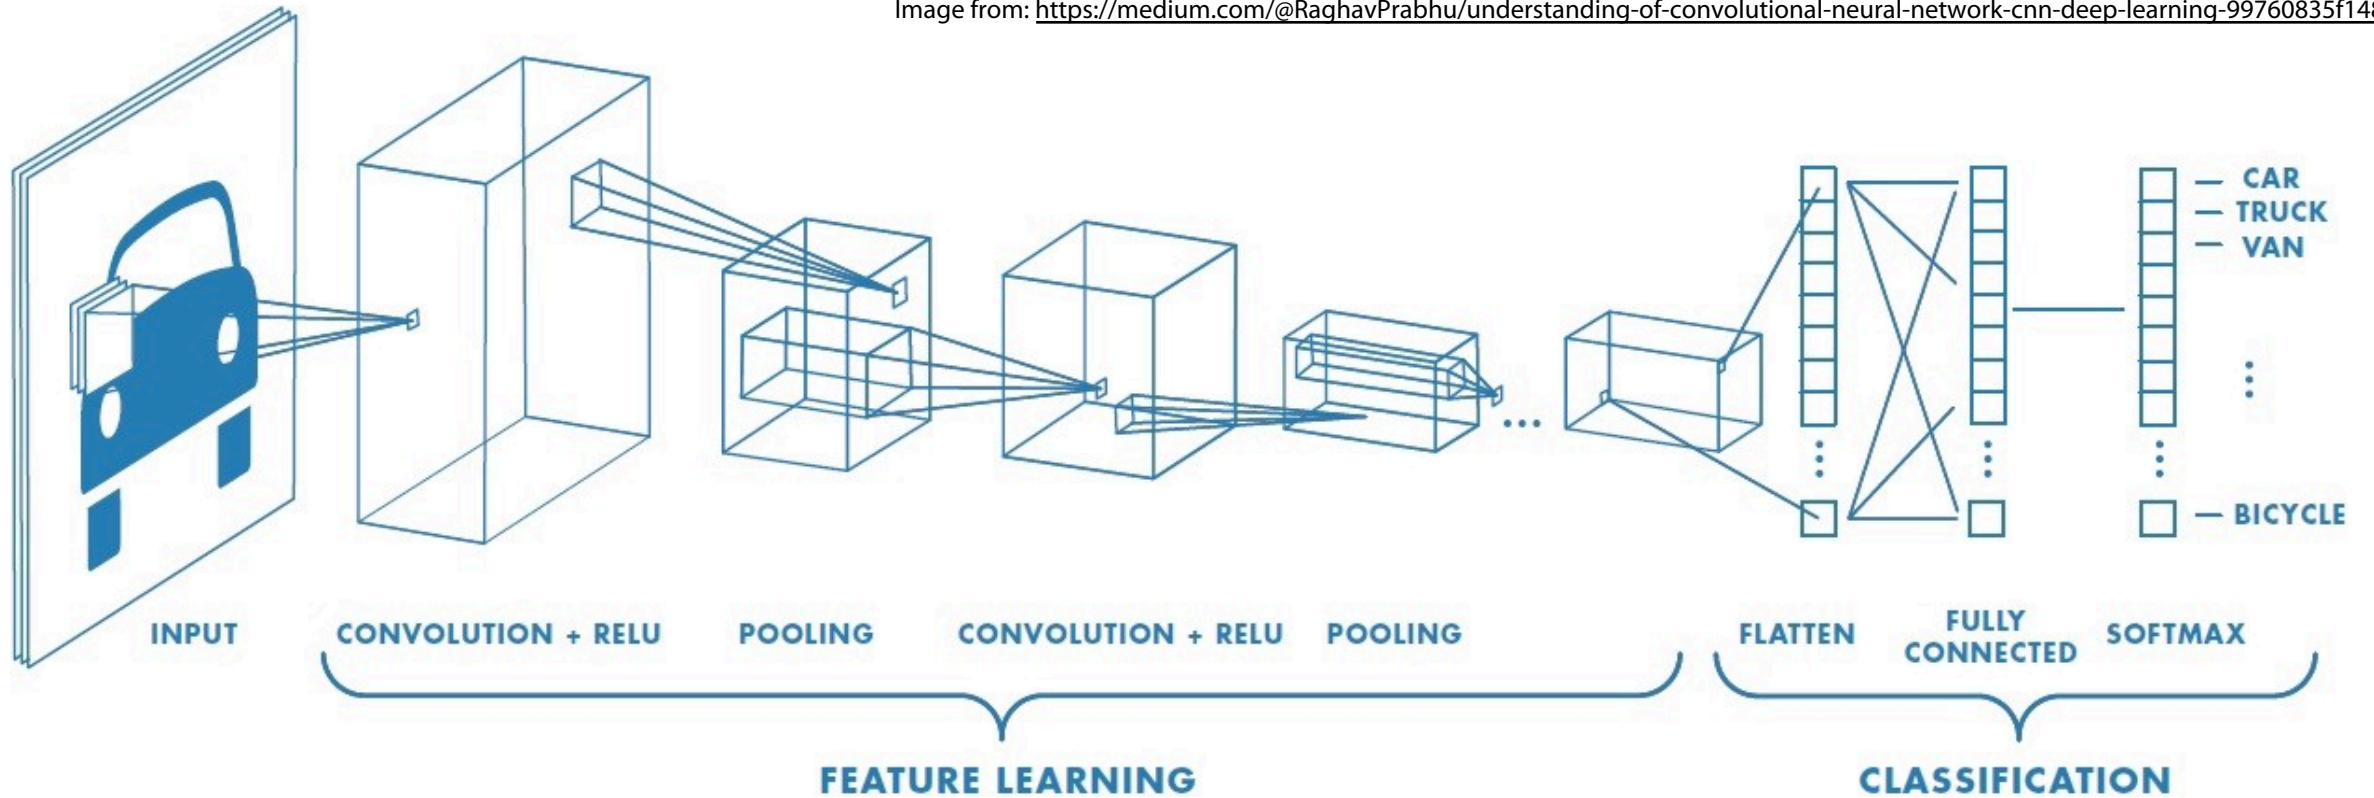
1	-2	5	7
2	1	4	-3
4	-9	2	0
-4	15	2	1

$\xrightarrow{\text{maxpool}(2 \times 2)}$

2	7
15	2

# Deep convolutional network

Image from: <https://medium.com/@RaghavPrabhu/understanding-of-convolutional-neural-network-cnn-deep-learning-99760835f148>



- Convolution is typically followed by a pooling operation (aggregating values of nearby pixels), e.g. maxpool:

## Motivation:

Deeper layers extract "higher level" (i.e. more complex) features, and we're less interested in their spacial location. Hence, fewer pixels and more channels.

4	5	2	0
-4	15	2	1

15	7
2	2

# Thank you!



[amaevskij@hse.ru](mailto:amaevskij@hse.ru)



SiLiKhon



hse\_lambda

Artem Maevskiy